

本科生《计算机视觉》 基于深度学习的视觉理解与生成

黄雷

人工智能研究院

huangleiAI@buaa.edu.cn

2023年10月12日

主要内容


- 深度学习基础
 - 神经网络及反向传播算法
 - 卷积神经网络中的视觉表示思想
- 视觉理解任务
 - 目标检测
 - 分割
- 视觉生成
 - 深度生成模型
 - 图像翻译任务详解
- 深度神经网络训练技巧

outline


- Modeling of CNN
 - Module-wise architecture 模块化结构
- Convolutional layer (module)
 - Convolution like template
 - Filters
 - Convolution module
- Pooling layer (module)

Feature extraction

- Feature extraction
 - Pixel-wise input
 - Correlation between features




Convolutional Neural
Network(CNN),卷积神
经网络




Convolution Neural Network

- Lenet-5



Convolution related layers

全连接层



outline

- Modeling of CNN
 - Module-wise architecture 模块化结构
- Convolutional layer (module)
 - Convolution like template
 - Filters
 - Convolution module
- Pooling layer (module)



Module-wise architecture

➤ PyTorch 平台

- **Model construction**

```
# Define model
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10)
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```



Module-wise architecture




➤ PyTorch 平台

```
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)
```

• Training per iteration:

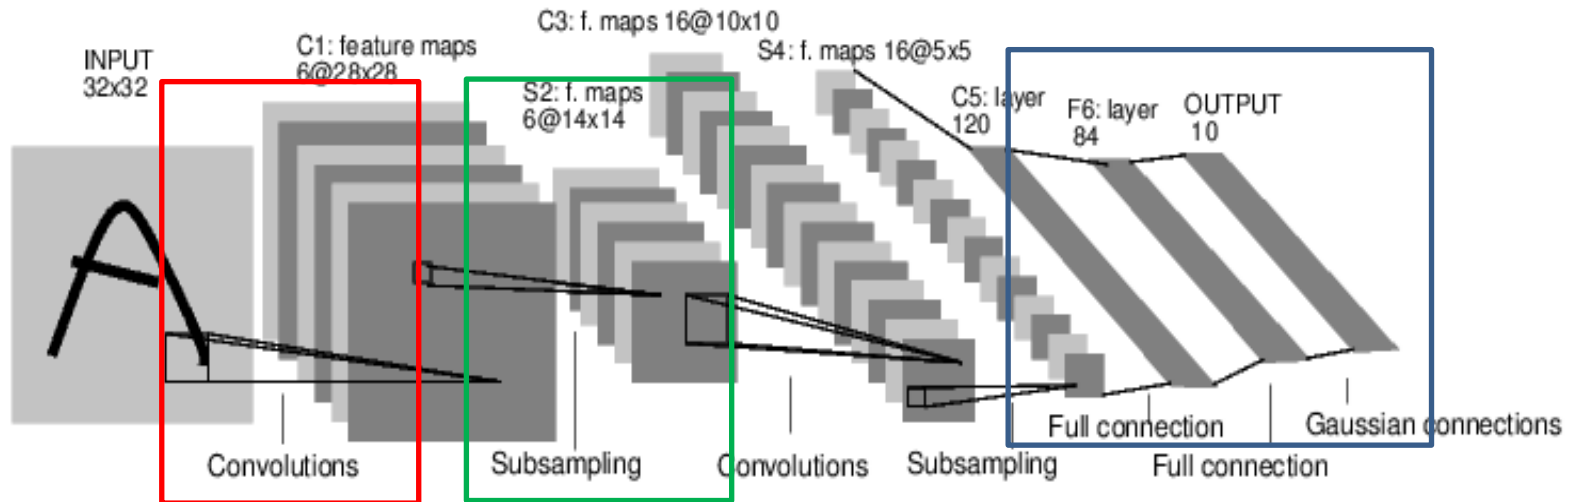
```
# Compute prediction error
pred = model(X)
loss = loss_fn(pred, y)

# Backpropagation
optimizer.zero_grad()
loss.backward()
optimizer.step()
```



Convolution Neural Network


- Lenet-5



Convolution (卷积)

Pooling (池化)

全连接层




outline

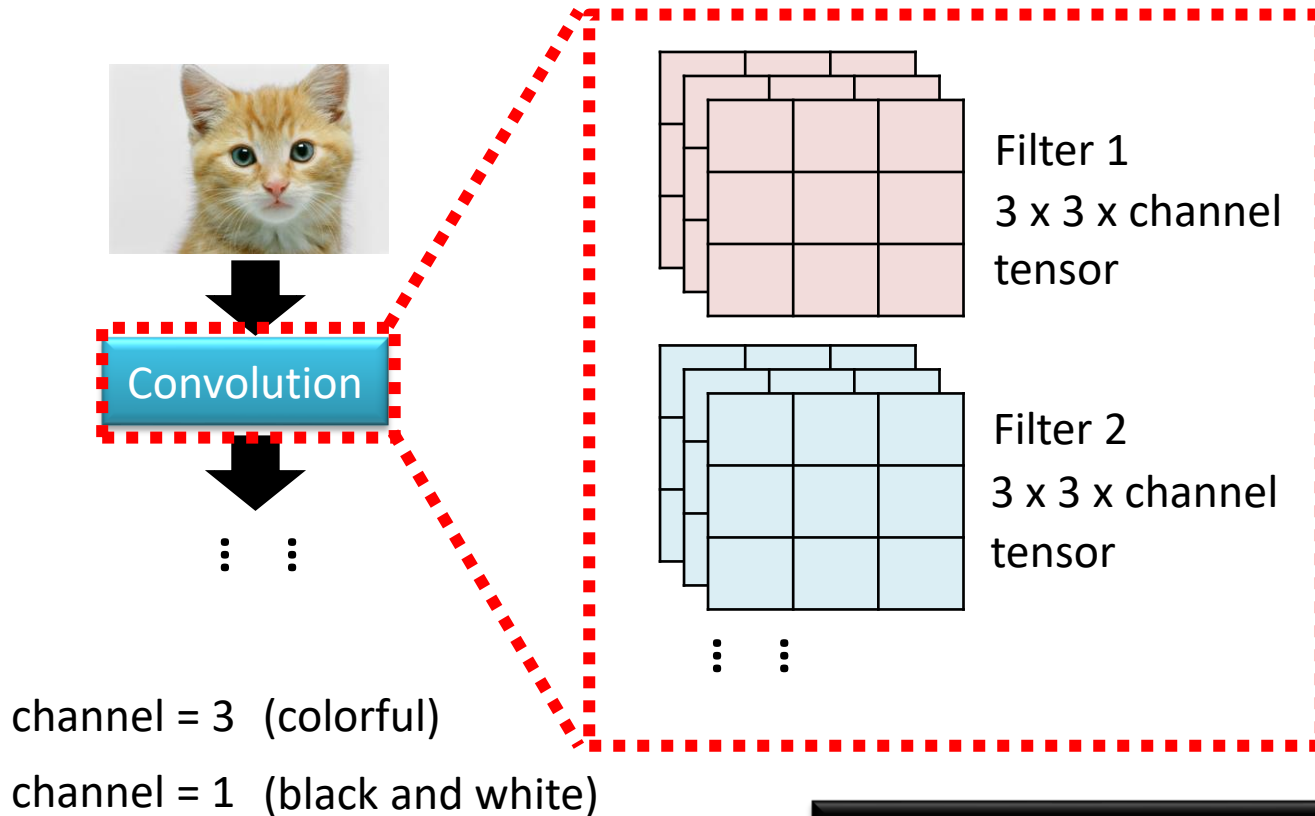
- Modeling of CNN
 - Module-wise architecture 模块化结构
- Convolutional layer (module)
 - Convolution like template
 - Filters
 - Why convolution for vision
- Pooling layer (module)

Convolution

- Match like template



Convolutional Layer



Each filter detects a small pattern (3 x 3 x channel).

Convolutional Layer

Consider channel = 1
(black and white image)

| | | | | | |
|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 1 | 0 | 0 |
| 1 | 0 | 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 1 | 0 |

6 x 6 image

| | | |
|----|----|----|
| 1 | -1 | -1 |
| -1 | 1 | -1 |
| -1 | -1 | 1 |

Filter 1


| | | |
|----|---|----|
| -1 | 1 | -1 |
| -1 | 1 | -1 |
| -1 | 1 | -1 |

Filter 2

⋮ ⋮

(The values in the filters
are unknown parameters.)

Convolutional Layer



Convolutional Layer

stride=1


| | | | | | |
|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 1 | 0 | 0 |
| 1 | 0 | 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 1 | 0 |

6 x 6 image

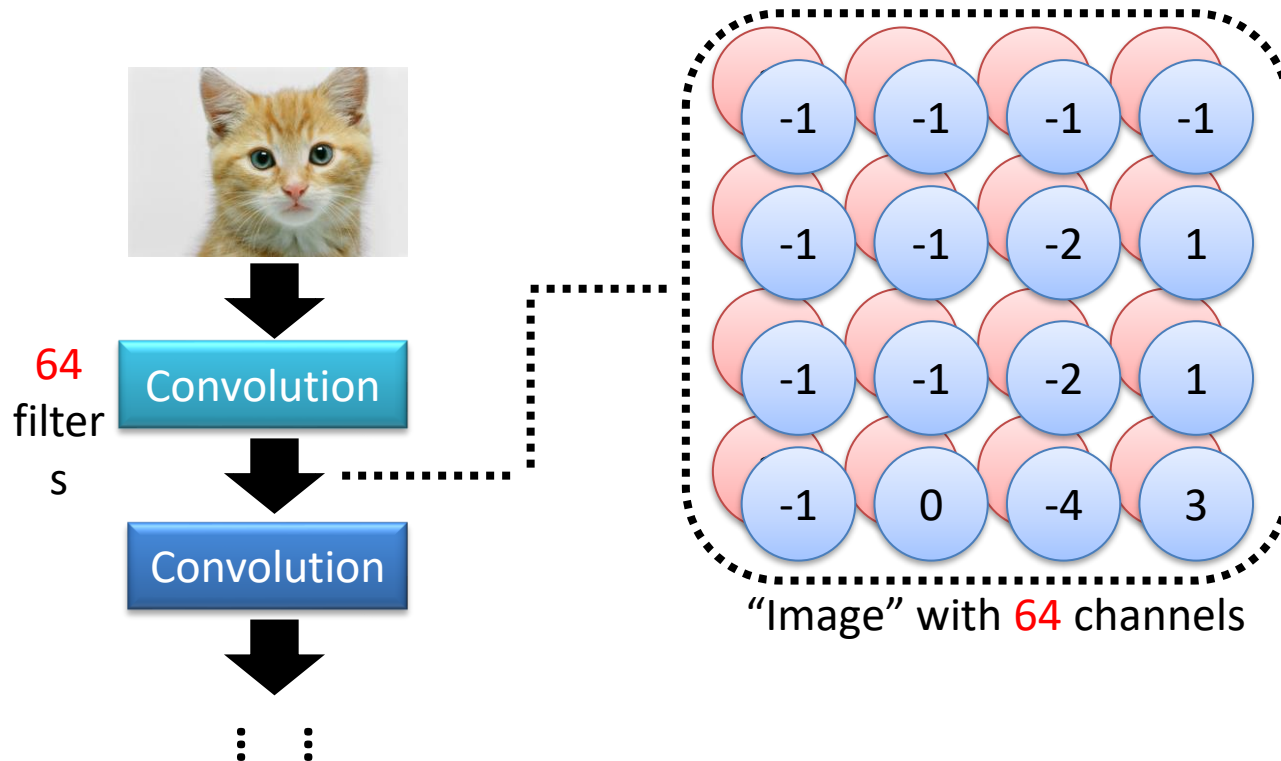
| | | |
|----|---|----|
| -1 | 1 | -1 |
| -1 | 1 | -1 |
| -1 | 1 | -1 |

Filter 2

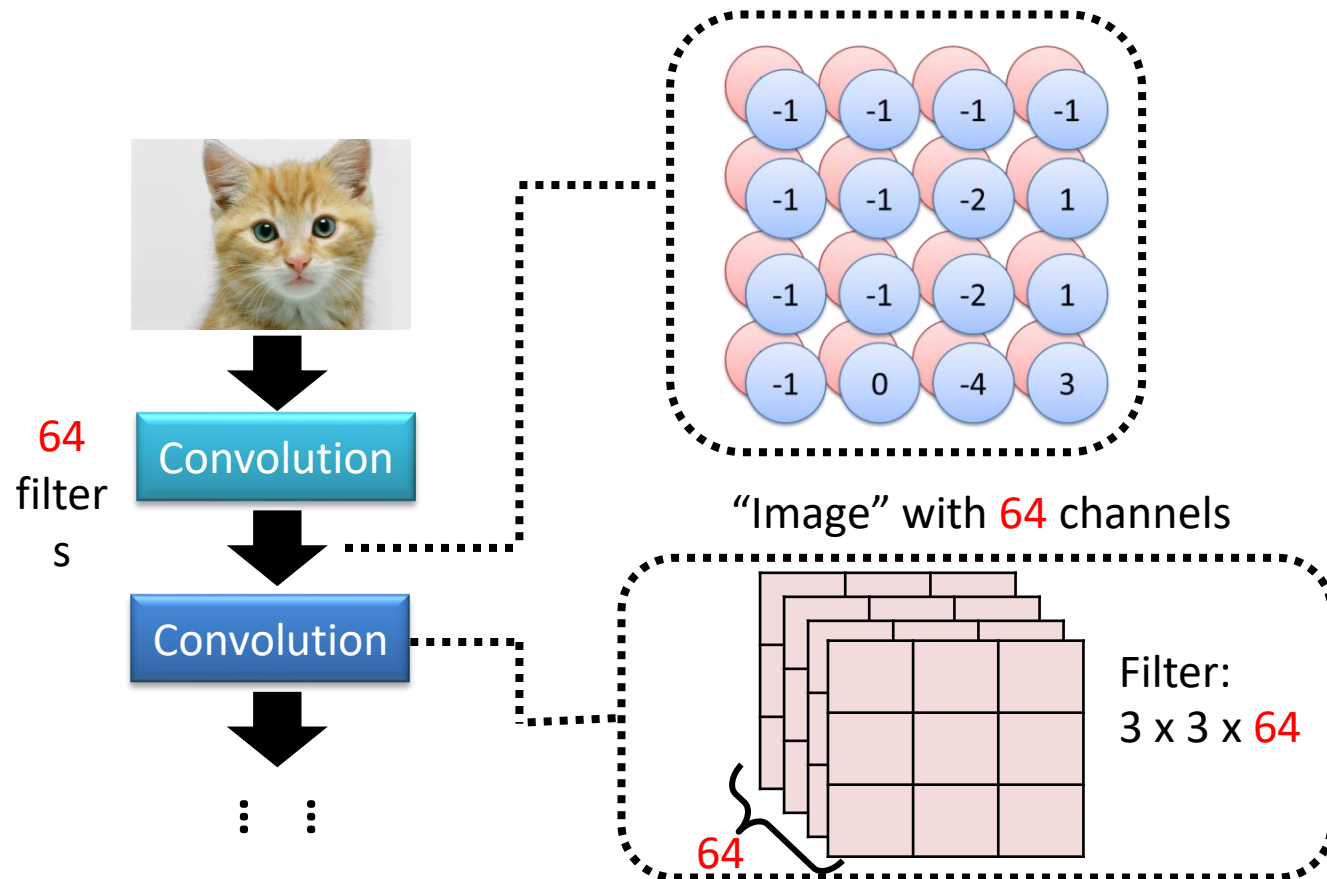
Do the same process for every filter




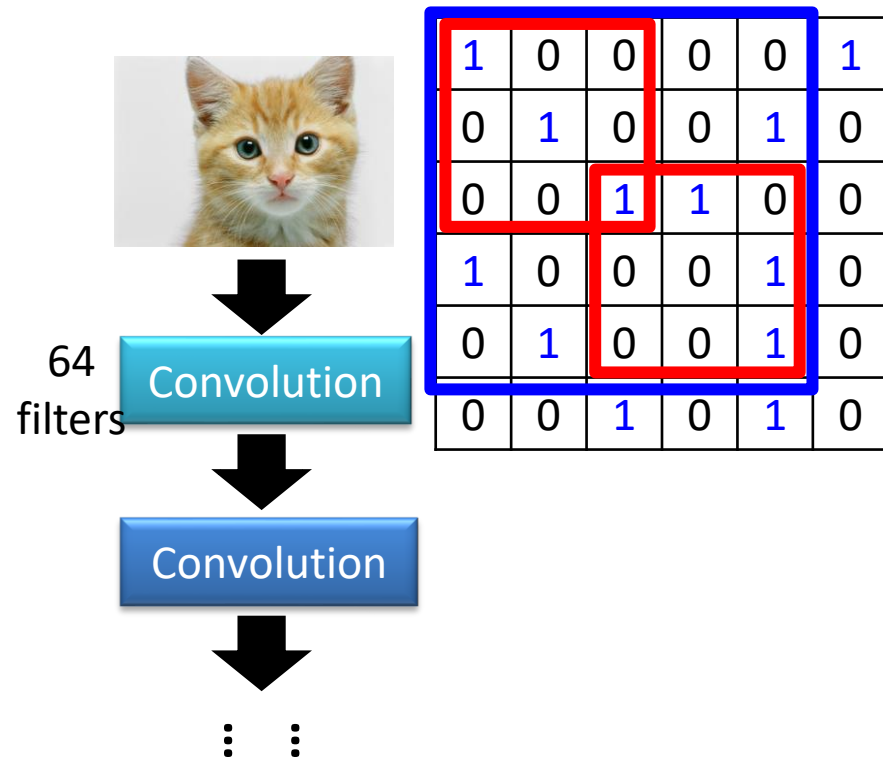
Convolutional Layer



Convolutional Layer



Receptive Field (感受野)




Convolution Layer (卷积层)

Input: $\mathbf{X} \in \mathbb{R}^{d_{in} \times h \times w}$

weight: $\mathbf{W} \in \mathbb{R}^{d_{out} \times d_{in} \times F_h \times F_w}$


output: $\mathbf{Y} \in \mathbb{R}^{d_{out} \times h \times w}$



Input: $\mathbf{x} \in \mathbb{R}^{d_{in}}$

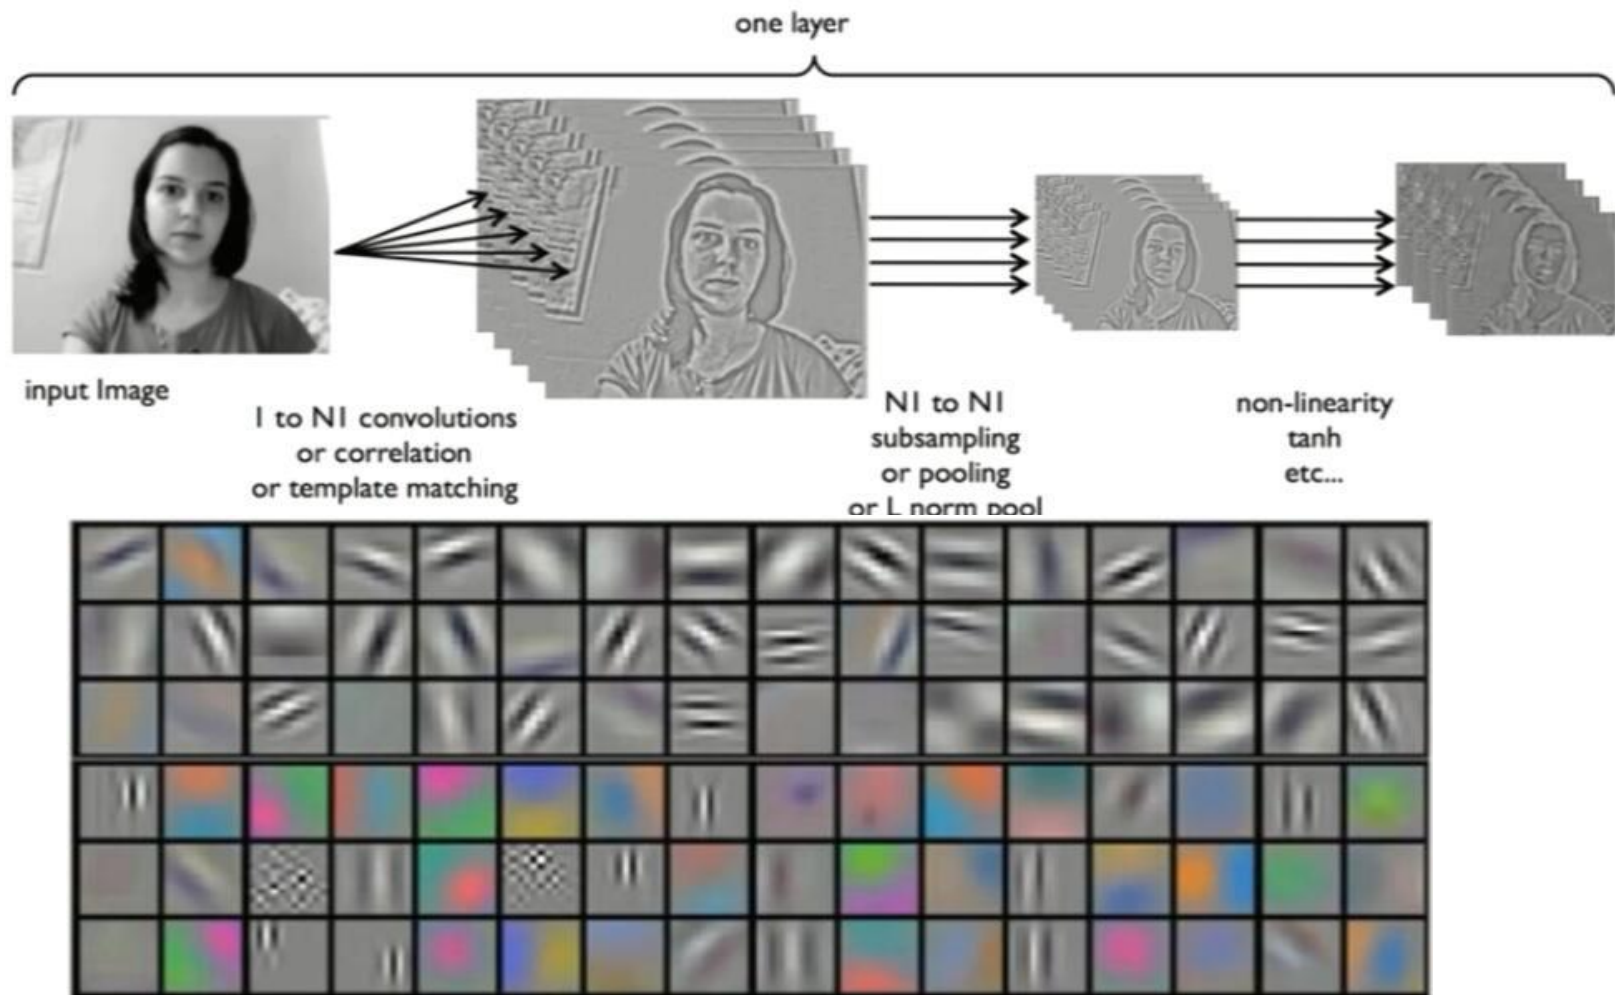
weight: $\mathbf{W} \in \mathbb{R}^{d_{out} \times d_{in}}$

output: $\mathbf{y} \in \mathbb{R}^{d_{out}}$



Feature detection (特征检测)

- Learning filters (weights)



outline

- Modeling of CNN
 - Module-wise architecture 模块化结构
- Convolutional layer (module)
 - Convolution like template
 - Filters
 - Why convolution for vision
- Pooling layer (module)

Practice with linear filters(线性滤波器)



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |

Filter



Filtered
(no change)

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 0 | 0 |

Filter



Shifted *left*
By 1 pixel

Practice with linear filters




Original

$$\frac{1}{9}$$

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

Filter



Blur (with a
box filter)

Practice with linear filters



Original

| | | |
|---|----|---|
| 0 | 1 | 0 |
| 1 | -4 | 1 |
| 0 | 1 | 0 |

Filter



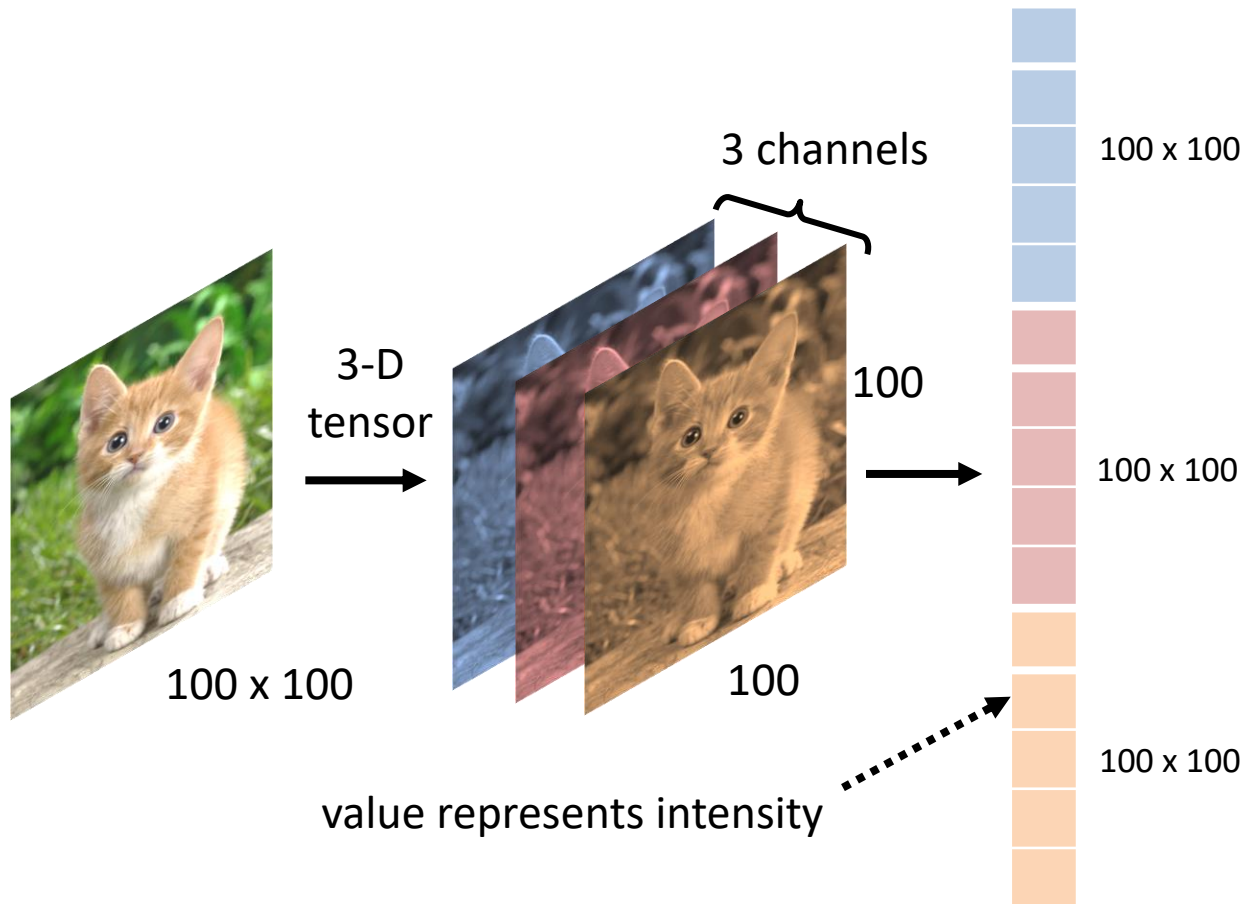
Output Image


Edge detect (边缘检测)

outline

- Modeling of CNN
 - Module-wise architecture 模块化结构
- Convolutional layer (module)
 - Convolution like template
 - Filters
 - Why convolution for vision
- Pooling layer (module)

Image Classification



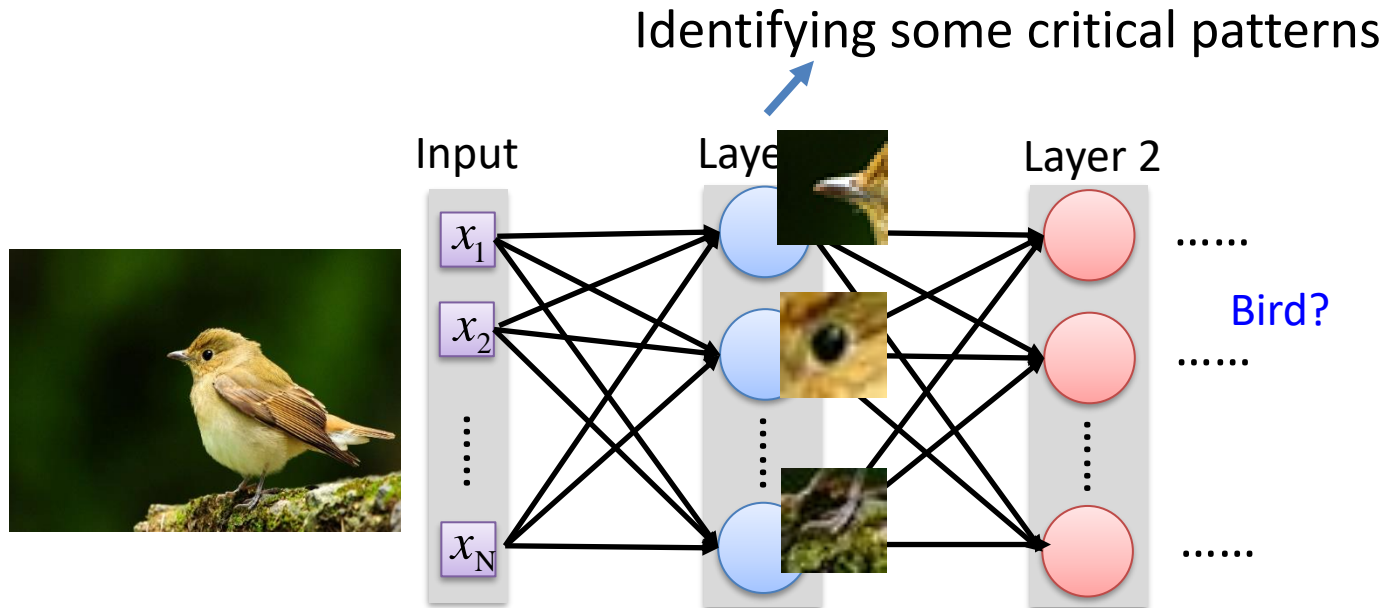


Vision model

How MLP to CNN?


- Model locality
- Parameter Efficiency

Observation 1




Observation 1

Need to see the whole image?




Some patterns are much smaller than the whole image.




A neuron does not have to see the whole image.

Simplification 1




Simplification 1



Simplification 1 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).



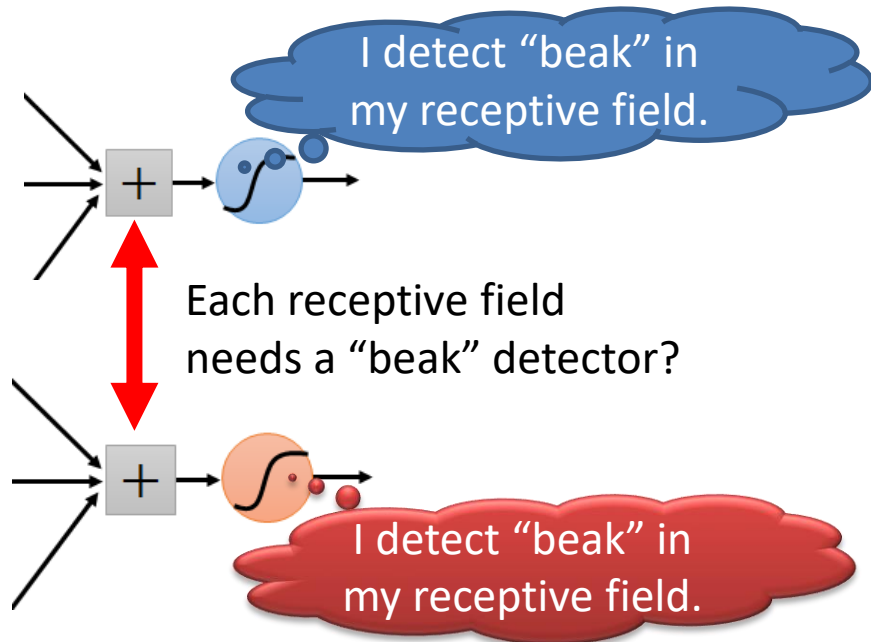
Vision model

How MLP to CNN?

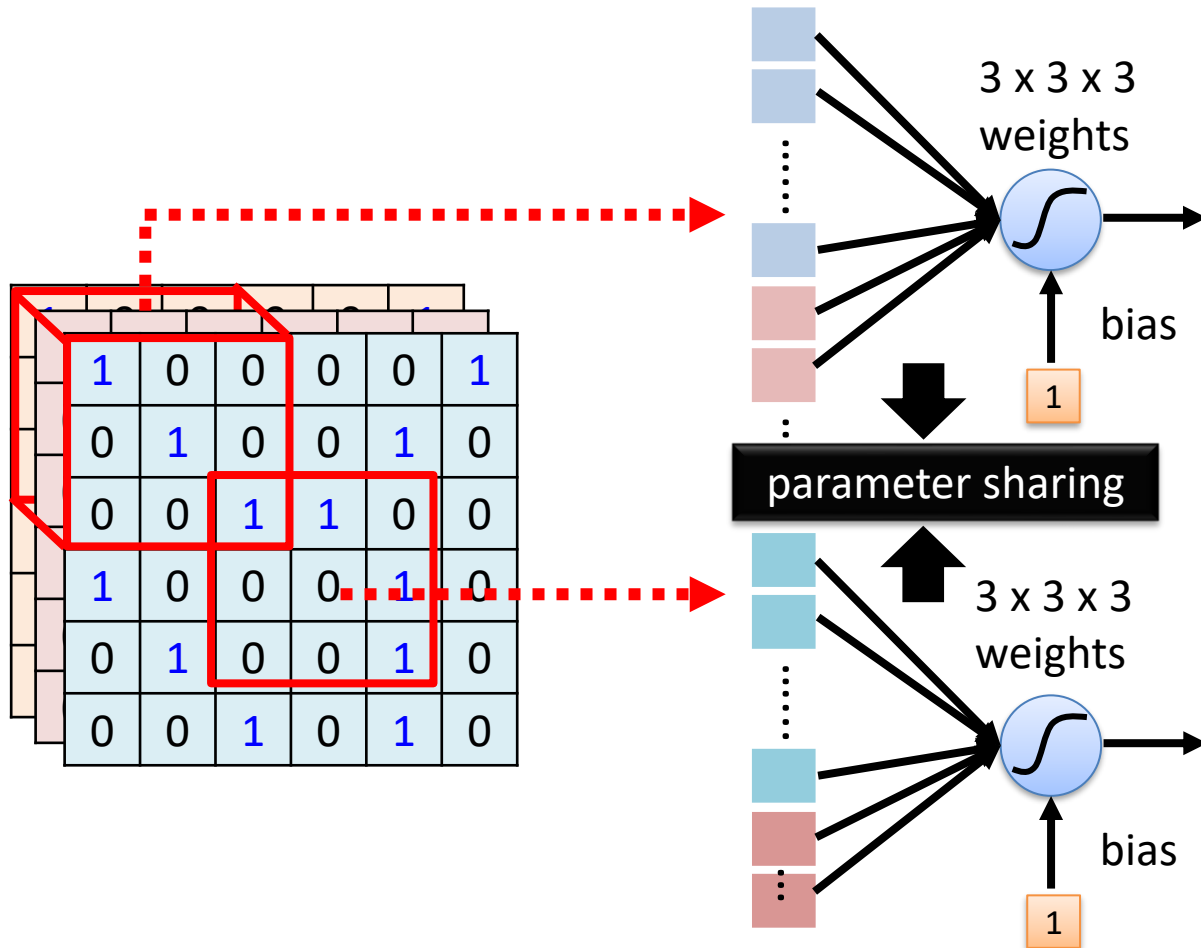
- Model locality
- Parameter Efficiency

Observation 2

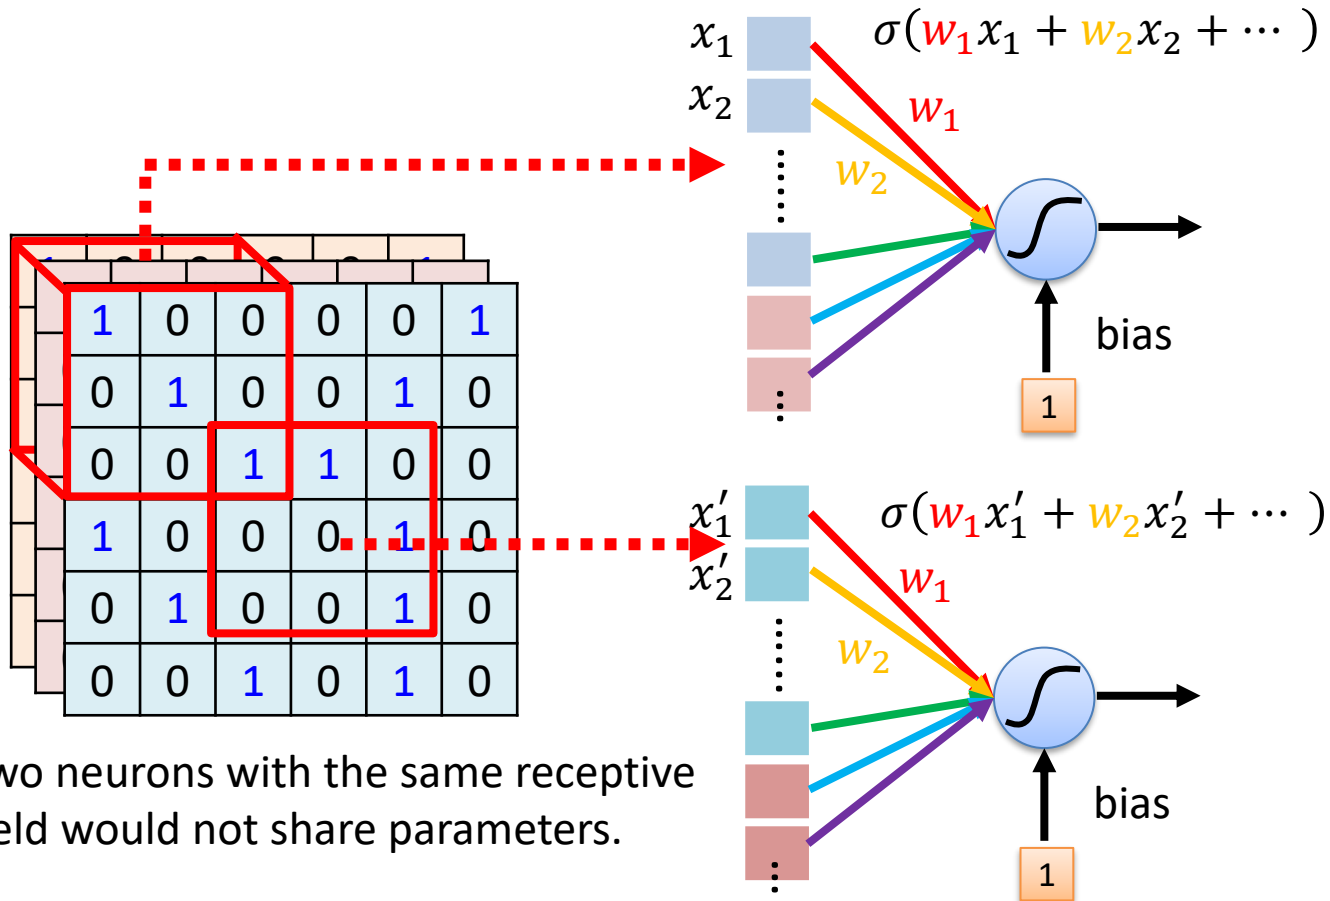
- The same patterns appear in different regions.



Simplification 2




Simplification 2



Simplification 2 – Typical Setting


Each receptive field has a set of neurons (e.g., 64 neurons).



Simplification 2 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).

Each receptive field has the neurons with the same set of parameters.




outline

- Modeling of CNN
 - Module-wise architecture 模块化结构
- Convolutional layer (module)
 - Convolution in general
 - Filters
 - Convolution module
- Pooling layer (module)


POOLING Layer

- In ConvNet architectures, **Conv** layers are often followed by **Pooling** layers
 - makes the representations smaller and more manageable without losing too much information.
 - Invariant in region.




MAX POOLING

Single depth slice




max pool with 2x2 filters
and stride 2




Average POOLING

Single depth slice




average pool with 2x2
filters and stride 2



Intuitive example


CONV CONV POOL CONV CONV POOL CONV CONV POOL FC
ReLU ReLU ReLU ReLU ReLU ReLU ReLU ReLU ReLU (Fully-connected)




Source: Andrej Karpathy & Fei-Fei Li

The characteristics of CNN

- Equivalent (等变性)----object detection
- Invariant (不变性)----Image Classification




receptive field (感受野)




Activation on Feature map

Deep neural networks visualization


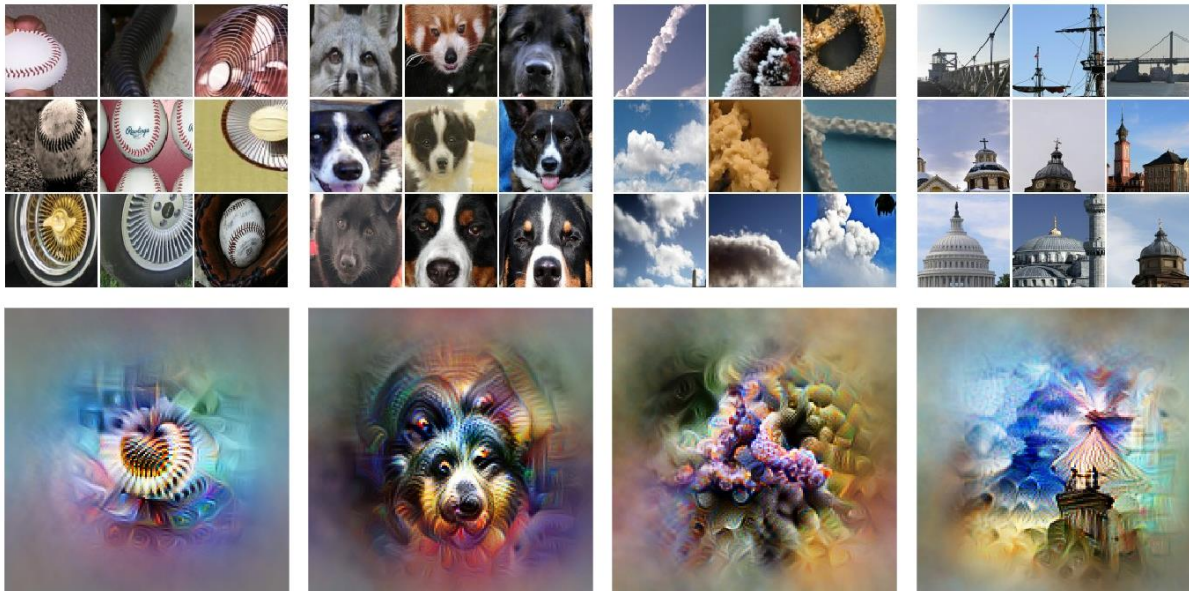

- 对抗样例 (Adversarial example)




$f(X + \Delta X) \neq f(X)$, ΔX should be imperceptible by human



Deep neural networks visualization



Neural Style Transfer




- Style
- Content



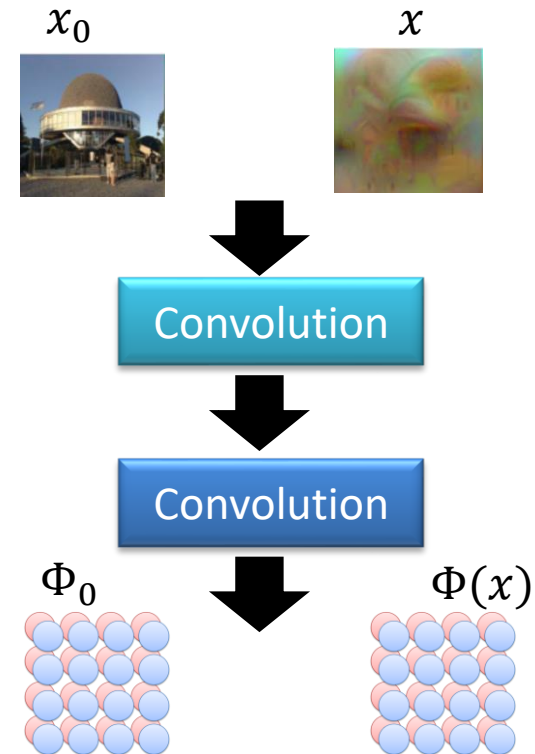
Reconstructing an image from a convolutional layer

- Representation function: $\Phi : \mathfrak{R}^{H \times W \times C} \rightarrow \mathfrak{R}^d$ (image space to feature space)
- Target Representation: $\Phi_0 = \Phi(x_0)$ (x_0 is the original image)
- We need to find: $x \in \mathfrak{R}^{H \times W \times C}$ by minimizing:

$$x^* = \arg \min_{x \in \mathfrak{R}^{H \times W \times C}} l(\Phi(x), \Phi_0) + \lambda R(x)$$



“Understanding Deep Image Representations by Inverting Them”, by Aravindh Mahendran and Andrea Vedaldi.




Content Loss Function

- Filters (Depths) at layer l : N_l
- The height times the width of the feature map at layer l : M_l
- Response at layer l : $F_l \in \mathbb{R}^{N_l \times M_l}$

F_{ij}^l represents the i th filter at position j in layer l

- Original image: \vec{p}
- We generate image: \vec{x} (randomly initialized)
- Squared-error loss:

$$L_{content} = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2$$



Style Loss Function

- Filter correlations are given by the Gram matrix:

$$G^l \in \mathbb{R}^{N_l \times N_l}$$

- G^l_{ij} is the inner product between the filters i and j in layer l :

$$G^l_{ij} = \sum_k F^l_{ik} F^l_{jk}$$

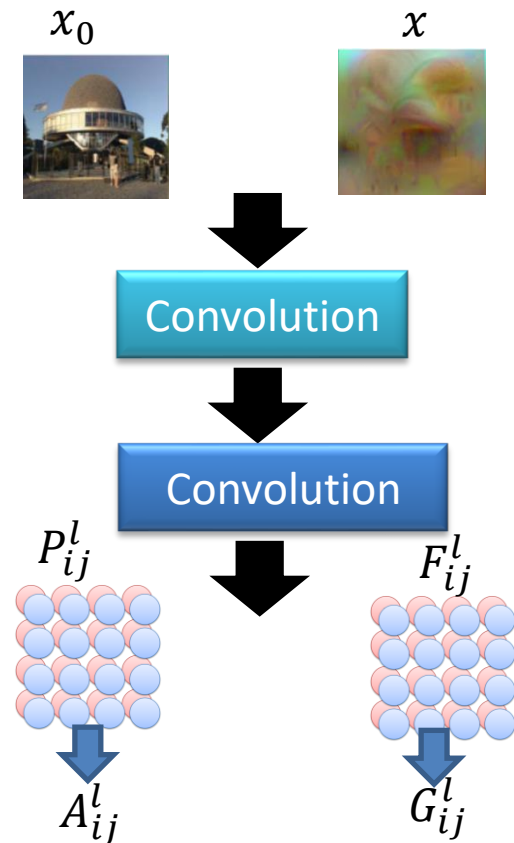
- The loss at layer l :

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G^l_{ij} - A^l_{ij})^2$$

A \leftrightarrow original image
G \leftrightarrow generated image

- The total style loss:


$$L_{style} = \sum_{l=0}^L w_l E_l$$



Content Reconstruction



Image reconstructed from layers
 (a)'conv1_1',
 (b)'conv2_1',
 (c)'conv3_1',
 (d)'conv4_1' and
 (e)'conv5_1'
 of the original VGG-Network




$$L_{content} = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2$$

Style Reconstruction



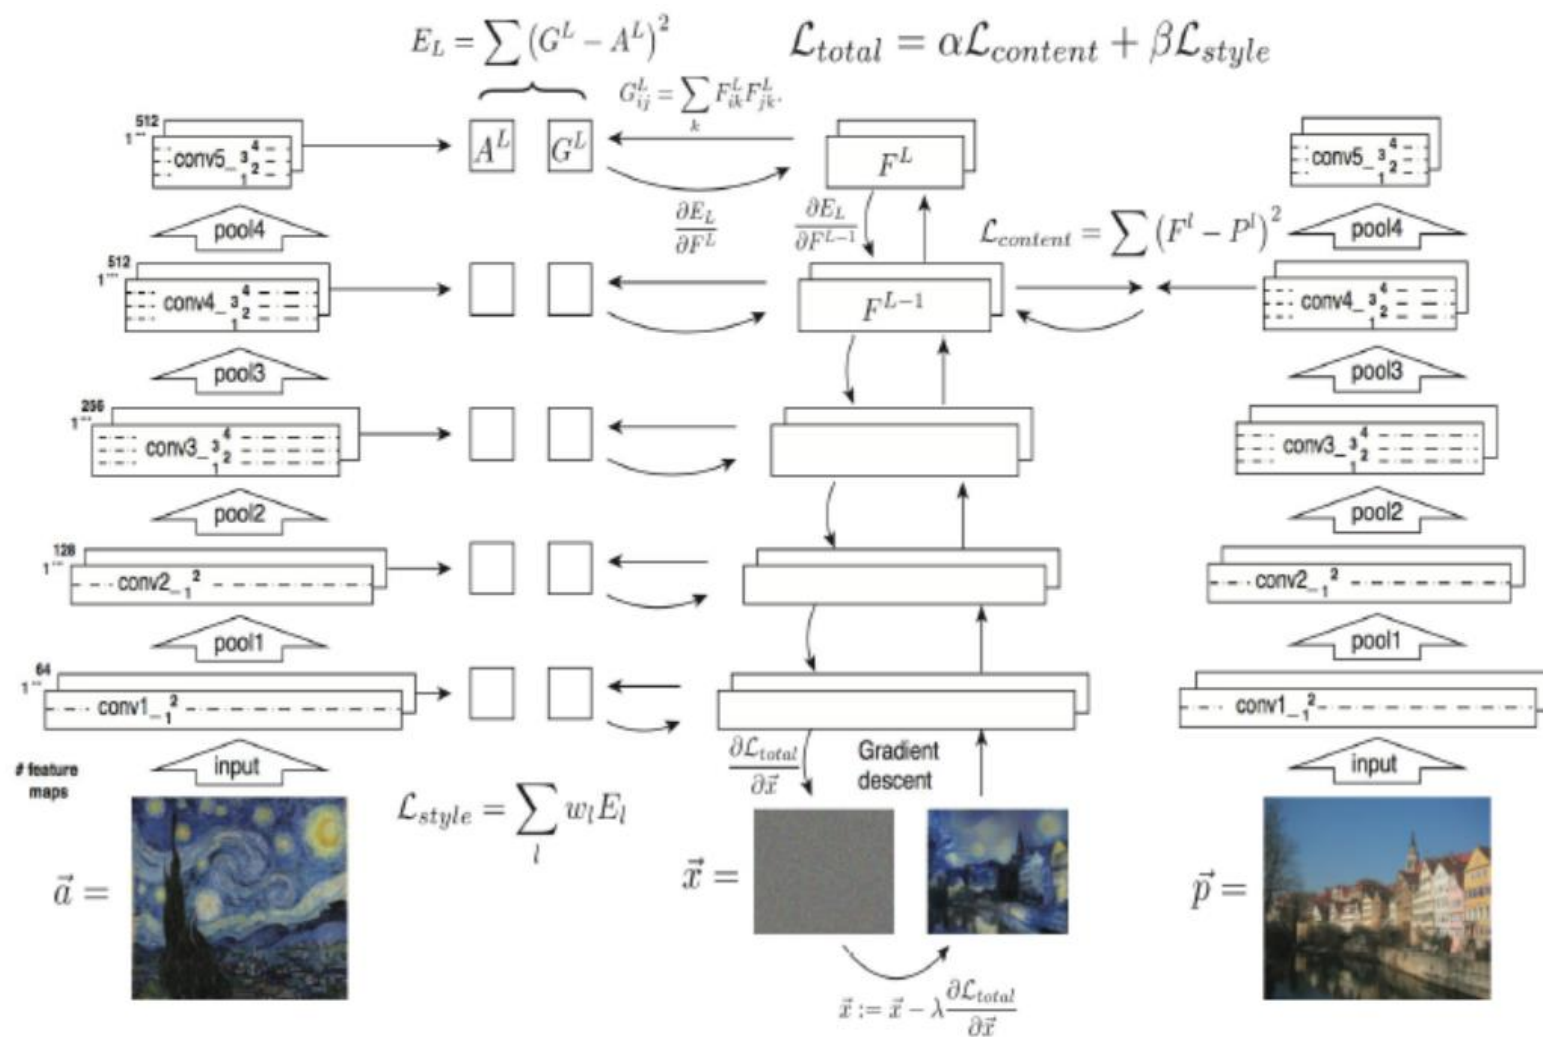
Style representations (filer correlations) from:

- (a) 'conv1_1',
- (b) 'conv1_1', 'conv2_1',
- (c) 'conv1_1', 'conv2_1', 'conv3_1',
- (d) 'conv1_1', 'conv2_1', 'conv3_1', 'conv4_1',
- (e) 'conv1_1', 'conv2_1', 'conv3_1', 'conv4_1', 'conv5_1'.




$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{ij}^l - A_{ij}^l)^2$$

The Total Loss Function



Results



Outputs at intervals
of a 100 iterations,
using white noise
for initialization

show image every 10 iterations



谢谢！

